The Talenti Decision System
– Edition April 2026 –

Decisioning Infrastructure.

Contents:

1. Decision Definition (Level 1 - Scope and Intent)
    a. What Decisions Are We Making?
    b. What Are The Outputs?
    c. What Does 'Good' Look Like?
    d. Objective
    e. Constraints
    f. System Principles
2. Decision Policy (level 2 - Rules)
    a. Behavioural Dimensions and Boundaries
    b. Signals
    c. Scoring
    d. Evidence Thresholds
    e. Conflict Resolution
    f. Risk and Failure Overlay
    g. Decision Stance Framework
    h. Environment Definition
    i. Mapping Engine
    j. Dimension Classification Framework
    k. Behavioural Priority Weighting
    l. Decision Framing Layer
3. Decision Routines (level 3 - Execution)
    a. Interview Structure
    b. Question Design Principles
    c. Probe System
    d. Probe Escalation Logic
    e. Dimension Coverage Rules
    f. Decision Tension Enforcement
    g. Evidence Depth Enforcement
    h. Signal Integrity Protection
    i. Time & Depth Constraints (MVP1)
    j. Interviewer Control Rules
    k. Scenario Library
4. Decision Architecture (level 4 - System Design)
    a. System components
    b. Data flow (input → signal <behaviour + skills> → score → decision)
    c. Determinism vs flexibility

    d. LLM usage and constraint boundaries
    e. Identity and data handling (Tokenisation process(?), Data separation layers, Mapping logic (ID ↔ real profile)
    f. Archetype Layer: Contextual Intelligence & Pattern Learning
    g. LLM Execution Layer: Prompting & Scoring Enforcement
5. Decision Governance (level 5 - Control)
    a. Override rules
    b. Audit triggers
    c. Failure state classification (system vs candidate)
    d. Drift detection (inc archetype drift detection)
    e. Feedback loop design
    f. Policy update rules - who can change scoring logic / what triggers a change?
    g. Decision audit cadence - what gets reviewed weekly/months?
    h. System performance metrics - what defines a good decision system? Inc archetype prediction accuracy
    i. Bias minimisation controls - Identity tokenisation enforced pre-interview, No identity exposure in scoring pipeline, Audit logs for data access
6. Cross-Cutting Controls: Signal Integrity (Applies across Policy, L2; Routines ,L3; and Architecture, L4) Priority 4
    a. Anti-gaming logic
    b. Consistency checks
    c. Depth enforcement
    d. Multi-angle validation
    e. Pattern detection (rehearsed vs real)
    f. Signal invalid rules - when do we say 'this does not count' (overly generic examples, hypothetical answers described as real, repeated frameworks/cookie cutter answers - consider tying to boundaries in 2c.)
    g. Skills signal validation rules
7. Decision Contract (level 6 - Output) Priority 2
    a. Output structure (Proceed / Proceed with Conditions / Do Not Proceed)
    b. Confidence - Behavioural Confidence + Skills Confidence (evidence-based)
    c. Integrity (clean / conflicted / insufficient)
    d. Explanation format
    e. Decision traceability - what evidence led to this decision / can it be replayed?
    f. Archetype alignment signal - incs. Candidate + role archetype classification + alignment confidence (eventually becoming visible)
    g. Candidate Feedback Layer

System Execution Flow:

1. Define environment (2h)
2. Map critical dimensions (2j)
3. Run interview + extract signals (3)
4. Validate signals (2c)
5. Check evidence thresholds (2e)
6. Score dimensions (2d)
7. Resolve conflicts (2f)
8. Apply decision thresholds (2g)

STEP	1 DEFINE ENVIRONMENT (2g) Environment Definition	2 MAP ENVIRONMENT TO DIMENSIONS (2h) Mapping Engine	3 CLASSIFY DIMENSIONS (2i) Dimension Classification Framework	4 SET PRIORITIES (2j) Behavioural Priority Weighting	5 DEFINE DECISION FRAMING (2k) Decision Framing Layer	6 ESTABLISH SIGNALS (2b) Signals	7 DEFINE SCORING (2c) Scoring	8 SET EVIDENCE THRESHOLDS (2d) Evidence Thresholds	9 DEFINE CONFLICT & INTEGRITY RULES (2e) Conflict Resolution	10 DECISION STANCE FRAMEWORK (2f) Decision Stance Framework	11 INTERVIEW HANDOFF (Section 3) Interview Preparation
ORG INPUTS (What the organisation provides)	• Role description • Operating context • Environment ratings: – Structure – Pace – Stakeholder complexity – Autonomy – Feedback	(No new input) Uses environment from Step 1	(No new input) Uses mapped thresholds from Step 2	• Top 1–3 behaviours they most need in this role/environment (e.g. Decision-making under ambiguity, Stakeholder navigation under pressure, Information gathering & clarification)	• Any role-specific dealbreakers • Additional notes/ context that should frame decisions (e.g. change tolerance required, travel, team leadership level)	(No input) Signals are defined by the system	(No input) Scoring rules are defined by the system	(No input) Threshold rules are defined by the system	(No input) Conflict types, resolution rules and integrity states defined by the system	(No input) Decision stance rules defined by the system	(No input) Interview is configured by the system (using all prior steps)
SYSTEM INPUTS (What the system uses/creates)	None (Starting point)	Environment ratings from Step 1	Mapped thresholds from Step 2	Priority behaviours from org (mapped to dimensions)	Role context and any dealbreakers from org	System library of validated behavioural signals (observed behaviours)	Dimension definitions and scoring scale (-2 to +2) and rules	Evidence density rules, confidence rules, minimum signals per dimension	Conflict rules, integrity status rules, tie-breaking rules, risk definition	Stance thresholds, decision logic, output rules	All outputs from steps 1–10
WHAT HAPPENS IN THE SYSTEM	Environment is captured and standardised by the system	Environment ratings are translated into dimension thresholds (+1 / 0 / N/A)	Dimensions are classified into: • Critical (≥ +1) • Minimum (≥ 0) • Non-critical (N/A)	Priority behaviours are mapped to dimensions and ranked (1–3)	Decision context and dealbreakers are recorded to frame evaluation	System establishes what behaviours count as evidence for each dimension (DDI filter applied: Direct, Distinctive, Observable, Lasting)	System defines how evidence translates to scores (-2 to +2) per dimension	System sets minimum evidence and confidence requirements per dimension (level-based)	System defines how conflicts are resolved and how integrity status is determined (Aligned, Mixed, At Risk, Invalid)	System applies all rules to define how final stance (Proceed / Proceed with Conditions / Do Not Proceed / Insufficient Evidence) is determined	Interview structure, questions and probes are configured (Section 3) **< Interview carried out >** Transcript generated → Input into system for analysis
KEY OUTPUTS FROM STEP	Standardised Environment Profile	Dimension Threshold Matrix	Dimension Classification (Table)	Priority Dimension List (Ranked 1–3)	Decision Framing Summary (Includes dealbreakers)	Signal Library (Per Dimension) What behaviours will be captured in interview	Scoring Guide (Per Dimension) -2 to +2 scale with definitions	Evidence Requirements (Table)	Conflict & Integrity Rules (Operational Guide)	Decision Stance Framework (Table of Rules & Outcomes)	
PURPOSE OF THIS STEP	Establish the operating context that defines what 'good' looks like in this role	Translate context into what levels of behaviour are required for each dimension	Define which dimensions are must-pass vs baseline vs supporting	Focus the decision on what matters most for success in this role	Apply role-specific lenses and boundaries before evidence is collected	Ensure only valid, observable behaviours are captured in interview	Create a consistent, objective way to score behavioural evidence	Ensure decisions are based on sufficient, reliable evidence	Manage conflict and evidence quality; ensure integrity of final decision	Define how all inputs combine to produce a clear, defensible decision	Prepare the interview to collect the evidence required by the system
FLOW SUMMARY	**Define Context (What matters)** →	**Map Requirements (What is needed)** →	**Classify Importance (What is required vs baseline)** →	**Set Focus (What matters most)** →	**Frame Decision (Apply role lenses & dealbreakers)** →	**Define Evidence (What counts)** →	**Define Meassrement (How it's scored)** →	**Set Standards (How much is enough)** →	**Handle Complexity (Conflicts & integrity)**	**Decide Outcome (What the result means)** →	**Prepare to Collect Evidence (Interview configured)**

OUTPUT OF STEP 11 → **Section 3 Interview is executed** → Transcript generated → Inputted into system for analysis (scoring, thresholds, conflict resolution, and final decision)

Diagram of workflow linked to section 2.

1. Decision Definition

1a. What Decisions Are We Making?

Talenti Decision System (TDS) makes **structured behavioural hiring decisions** based on **observed evidence**, not self-reported claims.

At its core, TDS determines:

- Behavioural execution fitness
  - → How likely is the candidate to operate effectively within the defined environment, based on observed behaviour
- Hiring decision
  - → Should the candidate **Proceed, Proceed with Conditions**, or **Do Not Proceed.**

TDS does not:

- Assess technical or functional skills (MVP1 constraint)
- Rank candidates against each other
- Optimise for “best overall candidate”
- Make subjective or impression-based decisions

It evaluates **fit for a specific role, in a specific environment, nothing else.**

1b. What Are The Outputs?

The system produces:

1. Behavioural dimension outputs

For each dimension:

- **Score**: -2 to +2 (strength and direction of evidence)
- **Confidence**: Low / Medium / High (reliability of evidence)

2. Decision outputs

- **Proceed** → Sufficient, consistent, reliable evidence
- **Proceed with Conditions** → Mixed, incomplete or conflicting evidence
- **Do Not Proceed** → Consistent negative or insufficient evidence
- **Insufficient Evidence (No Judgement)**

3. Supporting outputs

- **Evidence summary** (traceable)
- **Key strengths** (evidence-backed)
- **Identified risks** (behaviourally grounded)
- **Integrity status** (clean / conflicted / insufficient)

1c. What Does ‘Good’ Look Like?

A **good decision** is:

1. Evidence based

- Derived from observable behaviour and reasoning

2. Consistent

- Same inputs → same outputs
- Different candidates are assessed against the same standard

3. Traceable

Every decision can be explained:

- What was observed
- How it was scored
- Why it led to the outcome

4. Environment specific

- Judged relative to a defined operating context

5. Resistant to convincing factors

Not influenced by:

- Polished communication
- Rehearsed/isolated answers

6. Decision complete

The system either:

- Produces a defensible decision; OR
- Explicitly declares insufficient evidence

A **bad decision** is:

- impression-driven
- inconsistent
- unexplainable
- overly sensitive to weak signals
- forced despite insufficient evidence

1d. Objective

TDS is optimised to make defensible hiring decisions based on observed behavioural evidence of how a candidate operates within a defined environment.

TDS is not optimising for:

- Predicting potential
- Rewarding narrative quality
- Maximising candidate experience
- Increasing hiring volume

Core principle:

Decision over description. Evidence over claims. Context over generalisation.

1e. Constraints

These are our non-negotiable system limits:

1. Evidence constraint

No decision can be made without observable behavioural evidence.
→ Hypotheticals, opinions, or claims cannot drive scoring

2. Signal integrity constraint

Low-quality or invalid signals must be excluded
→ Generic, rehearsed, or unverifiable responses do not count

3. Minimum evidence constraint

Each dimension must meet a **minimum evidence threshold** before scoring
→ Otherwise: default = 0 + low confidence

4. Execution reality constraint

A candidate cannot proceed if they show:

- Consistent inability to deliver outcomes; or
- No evidence of completed work

5. Conflict constraint

Contradictory evidence within a dimension:
→ caps score at 0 and → triggers integrity = conflicted

6. No compensation constraint

Strength in one dimension cannot compensate for:

- Systemic weakness; or
- Failure in environment-critical behaviour

7. Environment binding constraint

All evaluation must be relative to the declared environment
→ behaviour that is “good” in isolation can still be wrong in context

8. Identity isolation constraint

No identity-linked data can influence any determining factor (signal, scoring, decisioning)

1f. System Principles

What the system does:

- **Evidence over claims**
  Decisions are based only on observed behavioural evidence, not self-described traits or intent
- **Relevance over narrative**
  Only signals directly tied to decision-making are considered; storytelling quality is ignored
- **Specificity over generalisation**
  Positive scoring requires specific, repeatable examples; not broad or abstract statements
- **Absence is valid**
  Lack of evidence is treated as a legitimate outcome, not a gap to be filled with inference
- **Context defines correctness**
  Behaviour is evaluated relative to the declared environment, not against generic standards
- **Consistency is enforced**
  Identical inputs must produce identical outputs, no discretionary interpretation

What the system **will not** do:

- **No performance bias**
  Articulation, confidence or polish cannot influence scoring
- **No hypothetical inference**
  Hypothetical, principle-based or unverified scenarios cannot drive positive scores
- **No relative ranking**
  Candidates are not compared against each other, only against defined requirements
- **No signal override**
  Strong isolated signals cannot compensate for systemic weakness or repeated failure patterns
- **No forced decisions**
  The system will not produce a decision when evidence is insufficient, low-quality or conflicting
- **No autonomous judgement**
  The system does not “decide” independently, outputs are the result of enforced rules and evidence constraints
- **No identity leakage**
  Identity-linked data is never used in signal extraction, scoring or decisioning
- **No proxy bias**
  Pedigree, background or indirect indicators of status cannot influence behavioural evaluation

2. Decision Policy

2a. Behavioural Dimensions and Boundaries

Objective:

Define each behavioural dimension with:

- Clear measurement intent
- Explicit boundaries
- Zero overlap in scoring

This section defines **what is being measured** (not how it is scored or validated).

Global rules (applies to all dimensions):

- Dimensions are independently observable
- No overlap in scoring intent
- Evaluation is based on observed behaviour only (via Signals)
- Environment is applied separately (2i / 2k)

Dimension 1 - Ownership

Measures:

Whether the individual **takes responsibility for outcomes beyond assigned tasks**

Includes:

- Initiating action without instruction
- Taking accountability for results
- Following through to completion
- Owning problems without deflection

Excludes:

- Execution quality
- Decision-making under uncertainty
- Stakeholder conflict

Scoring condition:

Must show clear personal accountability tied to a specific outcome.

If accountability cannot be isolated → **do not score**

Dimension 2 - Execution

Measures: ability to **convert intent into delivered outcomes under constraint.**

Includes:

- Task sequencing and planning
- Progress under pressure or limitation
- Completion and delivery
- Maintaining momentum

Excludes:

- Who initiated the work
- Handling ambiguity
- Stakeholder interaction

Scoring condition:

Execution requires: a completed or clearly progressed outcome

No outcome → **no execution score.**

Dimension 3 - Conflict and Challenge

Measures: Ability to **engage in constructive disagreement to improve outcomes**.

Includes:

- Challenging ideas or decisions
- Navigating resistance
- Influencing outcomes through disagreement
- Resolving tension productively

Excludes:

- General communication ability
- Ownership of outcomes
- Adaptability over time

Scoring condition:

Conflict requires: real interpersonal tension or opposition

No tension → **do not score.**

Dimension 4 - Ambiguity

Measures: Ability to **operate and make decisions under unclear, incomplete, or changing conditions**

Includes:

- Forming assumptions
- Making decisions without full clarity
- Adapting direction when information changes
- Creating structure where none exists

Excludes:

- Execution quality
- Ownership of outcomes
- Stakeholder conflict

Scoring condition:

Ambiguity requires: lack of clarity or structure at the point of action.

Clear instructions present → **do not score**

Dimension 5 - Feedback and Adaption

Measures: Ability to **update behaviour based on feedback and improve over time.**

Includes:

- Receiving feedback (positive or negative)
- Behavioural change in response
- Evidence of improvement
- Sustained change over time

Excludes:

- General openness to feedback
- Execution or delivery
- Conflict behaviour

Scoring condition:

Feedback requires: **demonstrated behavioural change over time**.

No change → **do not score**.

2b. Signals

Objective:

Define what constitutes a:

- **Valid behavioural signal**
- **Invalid or unusable signal**

Valid signal criteria:

A signal is only valid if it includes:

1. Real context

- Specific situation
- Real stakes
- Not generic or theoretical

2. Personal responsibility

- Candidate's direct role is clear
- Not team-level ambiguity

3. Observable action

- What they actually did
- Not what "was done"

4. Decision or behaviour

- A choice, judgement or behaviour
- Not just process description

5. Outcome

- What happened as a result
- Must be concrete

6. Causality

- Clear link: action → outcome
- Not correlation or assumption

If any of the above are missing:
→ Signal is **downgraded or invalid**

Invalid signal types:

These must **not** be scored:

a) Hypotheticals - “i would / i usually….”
→ Not real behaviour

b) Generic / framework answers - Textbook, templated, polished
→ Likely rehearsed or AI-assisted

c) Team blur - “We did…” without isolating contribution
→ No ownership clarity

d) Outcome-les stories - No result or impact
→ Cannot assess execution

e) Single-instance overreach - One example stretched across multiple dimensions
→ Cap scoring to +1

Signal strength classification:

Strong signal (+2 eligible)

- Specific, detailed
- Repeatable across contexts
- Includes constraint or difficulty
- Shows clear decision + outcome

Moderate signal (+1 eligible)

- One solid example
- Some detail
- Outcome present

Weak signal (-1 risk)

- Vague
- Partial structure
- Limited ownership or clarity

No signal (0)

- Talking ABOUT behaviour
- No actual example

Signal density rule:

High score requires multiple valid signals

Enforcement:

- 1 strong example → max +1
- 2 (or more) strong examples → +2 allowed

Cross-Dimension Signal Integrity Rule

One example cannot fully populate multiple dimensions.

Example:

- Candidate uses same story for:
    - Ownership
    - Execution
    - Ambiguity

→ Must extract **distinct evidence per dimension**

Otherwise:

→ cap scores
→ reduce confidence

Contradiction rule:

If signals within same dimension conflict:

- Score capped at 0
- Integrity = conflicted.

2c. Scoring

Objective:

Convert validated behavioural evidence into a **standardised, comparable score per dimension**

Scoring scale:

Score	Meaning
+2	Strong, repeatable, high-confidence evidence
+1	Solid, single-instance evidence
0	Not observed or insufficient evidence
-1	Weak, inconsistent, or concerning signal
-2	Clear negative pattern or failure

Scoring rules: (enforced by 2b + 2d)

+2 (strong positive)

Requires:

- 2+ strong signals
- Consistent behaviour across contexts
- Clear decisions + outcomes
- Meets evidence thresholds (2e)

+1 (positive)

Requires:

- 1 strong OR 2 moderate signals
- Clear action + outcome
- Meets minimum evidence

0 (neutral)

Used when:

- No valid signal
- Insufficient evidence
- Signals invalid per 2c

0 is not ‘average’, 0 is ‘unknown’.

-1 (weak/risk)

Assigned when:

- Behaviour is vague or inconsistent
- Partial ownership or weak execution
- Avoidance patterns or unclear outcomes

-2 (clear negative)

Assigned when:

- Clear failure behaviour
- Repeated negative signals
- Avoidance, deflection, or breakdown

Global scoring constraints:

1. No evidence → no score

If evidence thresholds (2e) not met: → default to **0**

2. No single signal inflation

One example: → max **+1**

3. No cross dimension reuse

Same example: → cannot justify multiple high scores
→ reduces confidence

4. Negative Signals Override Positives

If strong negative pattern exists: → cannot score positive

5. Confidence is separate

Score does not equal confidence:

- Score = what
- Confidence = how sure.

2d. Evidence Thresholds

Objective:

Define the **minimum quantity and quality of behavioural evidence** required before:

- Assigning a score
- Increasing confidence
- Progressing to decision

Core principle:

No score without sufficient evidence. No confidence without repeated evidence. Decisions cannot be finalised where confidence thresholds for critical or constrained dimensions are not met.

1. Minimum evidence requirement (per dimension)

Each behavioural dimension must meet baseline requirement (to allow any scoring at all):

- At least 1 valid signal (per 2b)
- That must include;
    - Clear action
    - Clear outcome
    - Clear ownership

If **not** met:

→ Score = **0**
→ Confidence = **Low**
→ Flag = **Insufficient Evidence**

2. Evidence depth requirement

To assign **+1**

Requires:

- **1 strong OR 2 moderate signals**
- At least:
    - one real decision point
    - one clear outcome

To assign **+2**

Requires:

- 2+ strong signals
- Across;
    - different contexts OR repeated behaviour
- Must show;
    - consistency
    - constraint or difficulty

If **not** met:

→ Cap score at **+1**

3. Evidence consistency requirement

Positive scoring requires consistency, not isolated performance

Rules:

- Signals must not contradict each other
- Behaviour must be **repeatable**, not one-off
- If inconsistency detected:
    → Cap score at **0**
    → Integrity = Conflicted

4. Signal density rule (cross dimensional)

One example does not equal full evaluation

Rules:

- A single example
    - → cannot support +2
    - → cannot fully populate multiple dimensions
- If reused across dimensions
    - → reduce confidence
    - → cap total scoring impact

5. Confidence assignment rules

Our confidence is NOT subjective; it's derived.

High confidence

- 2+ strong signals
- No contradictions
- Clear, detailed, consistent

Medium confidence

- 1 strong OR 2 moderate signals
- Minor gaps but usable

Low confidence

- Weak signals
- Missing structure
- Limited detail

6. Insufficient evidence handling

Insufficient Evidence is triggered IF ANY of:

- Critical dimension = 0
- Evidence thresholds not met (2b)
- Confidence below threshold

Outcome: → **Insufficient evidence (No Judgement)**.

7. Forced neutrality rule (lack of evidence must not be interpreted as positive)

Rules:

- No inference from tone, confidence or articulation
- No “benefit of the doubt” scoring

**Being deterministic and transcript based this won’t ever be an issue.

8. Evidence relevance rule (older or less relevant examples carry less weight)

Rules:

- Recent, role-relevant examples increase confidence
- Older, less relevant examples:
    - Do NOT cap score
    - BUT reduce confidence

9. Probe escalation trigger

Insufficient evidence must trigger deeper probing before scoring

If triggered:

- System must
    - → request additional examples
    - → increase specificity

If still unresolved: → default to **Insufficient evidence (No Judgement)** outcome.

→ Probe escalation and evidence completion logic are defined in 3.1 (Interview Design), which governs how insufficient signals are resolved prior to final decision evaluation.

2e. Conflict Resolution

Philosophy: Decisions are environment-relative, but constrained by non-negotiable behavioural thresholds. I.e., context matters but some things kill a candidate regardless.

Step 1 - Defining the non-negotiable constraint

EXECUTION Dimension = Primary constraint.

Execution operates as a global floor constraint across all environments. A violation of Execution (≤ -1 with sufficient confidence <Med-High Confidence>) restricts or blocks progression regardless of other signals.

Rule 1: Execution floor (material risk)

IF EXECUTION = **-1** (Medium+ confidence)
→ **CANNOT result in Proceed**

Outcome:

- **Proceed with Conditions** OR
- **Do Not Proceed**

Constraint:

Cannot be upgraded to **Proceed** via:

- Other dimensions
- Priority weighting
- Environment

Note: a result of **-1** with **LOW** confidence cannot be guaranteed, so low confidence rules apply.

Rule 2: Severe Execution failure (hard constraint)

IF EXECUTION = **-2** (Medium+ confidence)
→ **Do Not Proceed**

Outcome:

- **Immediate decision**
- No further evaluation required

Constraint:

- No override permitted

- No trade-off allowed
- Environment irrelevant

The overarching msg (validated by research, interviews, surveys and design partners); "We don't hire people who can't actually deliver". **We encode this.

Step 2 - Define Secondary Structure (contextual)

All other dimensions operate under contextual weighting (based on environment) but within guardrails.

Rule 3: Strength cannot mask systemic weakness

IF **Risk Stack ≥ 3 (with confidence ≥ Medium)** → CANNOT **Proceed**

Even if one dimension is +2 - this prevents 1 dimensional/broken candidates.

Rule 4: No Single Dimension Dominance

Max contribution of any single dimension is capped

I.e., one +2 does not carry the decision. Prevents over-indexing on 'star' signals.

Step 3 - Define Conflict Types

Conflict classified:

Type A: Direct Behavioural Conflict

Example:

- Claims Ownership (+2)
- But shows no follow-through (Execution -1)

Rule: Conflict between claimed behaviour and observed execution → downgrade both.

- Ownership capped at +1
- Execution remains dominant

Type B: Cross Dimension Tension

Tension between behavioural dimensions where strength in one creates risk in another, relative to the role environment. Common tensions:

1. Ownership vs Conflict & Challenge
2. Ownership vs Feedback / Adaptation

3. Execution vs Ambiguity
4. Conflict & Challenge vs Execution

Examples:

1. **Ownership (+) Vs Conflict and Challenge (-)**
    a. Pattern
        i. Takes responsibility, drives work (+ ownership)
        ii. Avoids pushing back, doesn't challenge (+ avoids friction)
    b. Risk
        i. doer, not challenger
        ii. will execute… but not correct bad direction
2. **Ownership (+) Vs Feedback and Adaption (-)**
    a. Pattern
        i. Strong agency, pushes forward
        ii. Low openness to feedback / low reflection
    b. Risk
        i. stubborn executor
        ii. scales mistakes fast
3. **Execution (+) Vs Ambiguity (-)**
    a. Pattern
        i. Can deliver in structured environments
        ii. Struggles without clarity
    b. Risk
        i. brittle operator
        ii. fails in dynamic environments
4. **Conflict and Challenge (+) Vs Execution (-)**
    a. Pattern
        i. Good at questioning / pushing back
        ii. Poor at actually delivering outcomes
    b. Risk
        i. smart talker / blocker
        ii. creates friction without progress

Scenario 1:

Signals:

- Ownership: +2
- Conflict & Challenge: -1
- Execution: +1

Environment:

- Highly autonomous, low debate culture

Outcome: **Proceed**

Reason:

- Challenge not critical
- Execution + ownership sufficient

Scenario 2: (with the same candidate)

Environment:

- High challenge / stakeholder-heavy role

Outcome: **Proceed with Conditions (if Execution ≥ threshold with Med+ confidence),** otherwise **Do Not Proceed.**

Reason:

- Avoidance of challenge = material risk
- Cannot operate effectively in environment

Scenario 3. (same candidate but **important** environment)

Signals:

- Conflict & Challenge: +2 (med confidence)
- Execution: -1 (med confidence)

Outcome: **Do Not Proceed** (global constraint override)

Reason:

- Execution floor violated
- Context irrelevant

Rules:

- Tensions are resolved relative to **declared environment requirements**
- **Global constraints override context** (e.g. Execution floor) however requires medium or high confidence to validate
- Weakness in **environment-critical dimensions cannot be offset**
- Strength in one dimension **does not compensate** for required capability gaps
- Execution constraints supersede all cross-dimensional and environmental considerations.
- Cross-dimensional strength may influence decision stance only where no global or critical constraints are violated.

Type C: Signal Density Conflict

Example:

- One strong example (+2)
- Everything else weak/absent

Rule: Low signal density → downgrade confidence AND cap score

- Max score = +1
- Confidence = Medium or Low

Type D: Contradiction Within Same Dimension

Example:

- First answer shows strong execution
- Later answer contradicts it

Rule: Intra-dimension contradiction → mark dimension as CONFLICTED

- Score capped at 0
- Integrity = conflicted

Following individual dimension evaluation, cross-dimension tensions are assessed to identify interaction risks prior to final decision resolution.

Step 4: Define Decision Outcome Logic

Proceed

Only if:

- Execution +1 (or more than)
- No dimensions -1 (or less than)
- Minimum 2 dimensions +1 (or more than)
- No critical conflicts
- Medium or High confidence

Proceed with Conditions

Triggered when:

- Execution 0 (or more than) but:
    - Conflicts present OR
    - Low signal density OR
    - 1 dimension -1 (or less than)

Guardrail as ‘not enough to trust’

Do Not Proceed

Triggered when:

- Execution -1 (or less than); OR
- Execution = 0 + multiple weak signals; OR
- 2+ dimensions -1 (or less than) ; OR
- Severe inconsistency (integrity failure)

Insufficient Evidence (No Judgement)

Triggered when:

- Not enough usable signals across dimensions

Step 5: Risks and Risk Accumulation Layer

Capture compounding risk across dimensions, not just isolated negative signals.

Risk definition:

A risk is recorded when:

- Score ≤ -1; AND
- Confidence ≥ Medium

Risk stack rules:

- **0–1 risks** → Acceptable (no escalation)
- **2 risks** → **Proceed with Conditions** required
- **3+ risks** → **Do Not Proceed**

Risk types:

- **Behavioural risk** → within a dimension
- **Structural risk** → environment misalignment
- **Integrity risk** → signal quality / conflict issues

Risk stack definition:

Count all dimensions where:

- Score ≤ -1
- Confidence ≥ Medium

Escalation conditions:

Escalate decision stance regardless of total stack if:

- **Any Priority Dimension ≤ -1 (Medium+ confidence)**; OR
- **Any Dimension = -2 (Medium+ confidence)**

Escalation outcome:

Cannot result in **Proceed**. Must result in:

- **Proceed with Conditions**; or
- **Do Not Proceed**

Step 6: Decision Integration Layer

Objective:

Translate evaluated signals into a **clear, accountable decision stance** with full explainability.

Definitions (explicit):

- **Decision** = the system's recommended course of action
- **Confidence** = strength of evidence supporting signals
- **Integrity** = stability and internal consistency of the decision given conflicting signals or gaps. Integrity does not change or override the decision outcome, but indicates the reliability and stability of that outcome..

Decision Structure:

Every output must follow:

Decision:

[Proceed | Proceed with Conditions | Do Not Proceed | Insufficient Evidence (No Judgement)]

Confidence:

- Low / Medium / High

Integrity definitions:

- **High Integrity** → signals consistent, low conflict
    - No intra-dimension contradictions
    - No cross-dimension unresolved conflict
    - Evidence thresholds met across all evaluated dimensions
    - Decision reached without forcing thresholds

- **Mixed** → some conflicting signals or gaps, but decision valid
    - Cross-dimension tension present (Type B conflict)

    - BUT decision still valid under rules
    - No critical constraint violations
    - Evidence sufficient

- **At Risk** → decision depends on trade-offs or unresolved uncertainty
    - Decision made under one or more of:
        - Low signal density across key dimensions
        - Borderline thresholds (e.g. relying on +1 with medium confidence)
        - Priority weighting materially influencing outcome
    - No hard invalidation, but **fragile decision**

- **Invalid** → insufficient evidence
    - Any of:
        - Critical dimension = 0
        - Insufficient evidence triggered
        - Signals fail validity rules (2b)
        - Decision would violate constraints if forced

Example:

Decision: Proceed with Conditions
Confidence: Medium
Integrity: Mixed

Reason:

- Strong ownership and collaboration signals
- Execution inconsistent across examples (Medium confidence)

Risks:

- Execution reliability under pressure
- Limited repeatable delivery evidence

Trade-off:

- Proceeding due to strong collaboration **despite execution risk**

>Decision = Stance + Confidence + Integrity + Trade-off<

= ENTERPRISE LEVEL EXPLANABILITY.

System risks:

- Execution misclassification - If signal detection is wrong, decision integrity collapses
  → Dependency: Level 3 (probing & extraction quality)
- Overuse of **Proceed with Conditions** - indicates weak thresholds/risk tolerance drift
  → Monitor distribution across decision outcomes

- Environment definition risk - Incorrect enviro inputs = systematically wrong decisions → Requires ongoing monitoring and validation (2i / 2k)

2f Risk and Failure Overlay

Objective:

Identify **compounding failure risk across dimensions within the defined environment**

Inputs:

- All -2 signals
- Any repeated -1 signals
- Stakeholder pressure behaviours (priority lens)

Environment-linked failure indicators

(derived from dimensions, not separate)

- Concedes under stakeholder pressure
- Avoids or defers conflict
- Delays decisions under ambiguity
- Fails to re-prioritise under constraint
- Breaks execution under load

Decision rule:

If:

- 2+ **critical failure indicators** present
  → **Decision = Do Not Progress**

If:

- 1 critical failure + weak supporting signals
  → **Decision = At Risk**

2g. Decision Stance Framework

Objective:

Translate scored behavioural dimensions into a final decision:

1. **Proceed**
    a. Clear, low-risk decision
2. **Proceed with Conditions**
    a. Decision made **with known risks and controls**
3. **Do Not Proceed**
    a. Decision made **based on sufficient negative evidence**
4. **Insufficient evidence (No Judgement)**
    a. No valid decision can be made

Inputs:

- Dimension scores (-2 → +2)
- Confidence (Low / Medium / High)
- Integrity status (2e)
- Critical + Minimum Dimensions (2k)
- Priority Dimensions (2k)

Decision States:

1. Proceed

Candidate shows:

- No unbounded compensation
- No conflicting confidence gates decisions
- No global constraints override context (Execution dimension) failed
- Keeping decisions must be explainable

2. Proceed with Conditions

Candidate shows:

- Identified risks; BUT
- Risks are manageable and understood

Trigger when:

- Any dimension ≤ -1 with Medium+ confidence; OR
- Risk Stack ≥ 2; OR
- Environment misalignment present

3. Do Not Proceed

Candidate shows:

- High-confidence risk; OR
- Compounded risk

Trigger when:

- Any dimension = -2 with Medium+ confidence; OR
- Risk Stack triggered (from 2e)
- Execution floor violated (from 2e)

4. Insufficient Evidence (No Judgement)

Trigger if:

- Any critical dimension lacks valid evidence
- Confidence = Low across key dimensions
- Signal density insufficient

Priority application:

Applied after:

- Critical thresholds satisfied
- Risk stack calculated

Applied only when:

- Decision = Proceed with Conditions OR borderline

Cannot:

- Override critical failures
- Reduce risk stack
- Upgrade Do Not Proceed → Proceed

Priority dimension sensitivity:

Rule: Priority dimensions adjust decision sensitivity across all non-terminal decision states.

5A - Positive reinforcement (not just Proceed with Conditions → Proceed)

If:

- Priority dimension ≥ +1
- Confidence ≥ Medium
- No blocking risks (risk stack < 3, no -2)

- All critical + minimum thresholds satisfied

→ **Upgrade decision stance by one level** (**Proceed with Conditions** → **Proceed**)

5B - Negative Escalation

If:
- Priority dimension ≤ -1
- Confidence ≥ Medium

→ Must degrade decision stance by one level

- **Proceed** → **Proceed with Conditions**
- **Proceed with Conditions** → **Do Not Proceed**

5C - Neutral

If:
- Priority dimensions = 0; OR
- Low confidence

→ **no change to stance**.

Constraints:

Cannot override:

- Execution floor (2e)
- Risk stack ≥ 3
- Any dimension = -2 (Medium+ confidence)

Cannot create progression from:

- **Do Not Proceed** → **Proceed**

Confidence X Severity rules:

- **-1 + Low confidence** → weak signal, does not independently influence stance but contributes to uncertainty
- **-1 + Medium confidence** → contributes to risk
- **-1 + High confidence** → must influence stance
- **-2 + Medium+ confidence** → cannot Proceed

Low-confidence signals increase uncertainty but do not act as standalone decision drivers. Any **priority dimension ≤ -1 with Medium confidence MUST influence decision stance**.

Environment misalignment rules:

If a candidate fails any **environment-critical dimension**. Then cannot be Proceed. It must be either:

- **Proceed with Conditions**; OR
- **Do Not Proceed**

Trade-off enforcement:

Every decision must include an **explicit trade-off statement**.

Format:

- Selected because: [Top strengths]
- Despite risk in: [Explicit risks]

Enforcement:

If risk(s) exist:

- The system **must not output a decision without stating a trade-off**.

System rules:

- No unbounded compensation across dimensions
- Confidence gates decision validity
- Global constraints override contextual factors
- All decisions must be explainable and auditable
- Insufficient evidence states take precedence over decision resolution and prevent final decision assignment.

2h. Environment Definition

Objective:

Define the operating context that **modifies how behavioural evidence is interpreted, scored and decided.**

Environment dimensions:

Each role must define:

Structure
- High (1) / Medium (2) / Low (3)

Pace
- Stable (1) / Moderate (2) / Fast (3)

Stakeholder Complexity
- Low (1) / Medium (2) / High (3)

Autonomy
- Low (1) / Medium (2) / High (3)

Feedback Culture
- Low (1) / Medium (2) / High (3)

Rules: Environment is **not** descriptive, **it is evaluative**.

1. Signal Interpretation.

The same behaviour can score differently depending on the environment.

Example:

- Low autonomy role → independent decision-making may be **misaligned**
- High autonomy role → same behaviour **= positive signal**

2. Threshold adjustment

Each dimension has **environment-sensitive thresholds**

Example:

- High ambiguity environment → Ambiguity must be ≥ +1 to **Proceed**
- Low ambiguity environment → Ambiguity may be non-critical

3. Failure conditions

Environment defines **which dimensions become critical**

Example:

- High stakeholder complexity:
  → Conflict & Challenge becomes **must-pass**
- Fast pace + low structure:
  → Execution + Ambiguity become **non-negotiable**

4. Misalignment rule

A behaviour that is correct in isolation but misaligned with the environment **must be scored as incorrect**.

5. Weighting

Environment determines:

- Which dimensions matter more
- Which weaknesses are tolerable
  BUT
- Does NOT allow compensation for hard constraints (from Section 1)

6. Default handling

If environment is:

- Missing → inferred with **low confidence flag**
- Vague → treated as **risk to decision integrity**

2i. Mapping Engine

Objective:

Translate environment inputs into:

- Critical dimensions (must-pass)
    - Execution is excluded from environment-based prioritisation and is always treated as a global constraint.
- Minimum thresholds per dimension
- Environment-driven failure conditions

1. Input

Environment profile:

- Structure (1–3)
- Pace (1–3)
- Stakeholder Complexity (1–3)
- Autonomy (1–3)
- Feedback Culture (1–3)

2. Dimension mapping rules

Autonomy → Ownership + Ambiguity

- HIGH (3)
  → Ownership ≥ +1 REQUIRED
  → Ambiguity ≥ +1 REQUIRED
- MEDIUM (2)
  → At least one of (Ownership, Ambiguity) ≥ +1
- LOW (1)
  → Neither required

Pace → Execution

- FAST (3)
  → Execution ≥ +1 REQUIRED
  → Execution -1 = **Do Not Proceed**
- MODERATE (2)
  → Execution ≥ 0
- STABLE (1)
  → Execution not required to exceed 0

Stakeholder Complexity → Conflict & Challenge

- HIGH (3)

→ Conflict ≥ +1 REQUIRED
- MEDIUM (2)
  → Conflict ≥ 0
- LOW (1)
  → Not required

Structure → Ambiguity

- LOW structure (3)
  → Ambiguity ≥ +1 REQUIRED
- MEDIUM (2)
  → Ambiguity ≥ 0
- HIGH (1)
  → Not required

Feedback Culture → Feedback & Adaptation

- HIGH (3)
  → Feedback ≥ +1 REQUIRED
- MEDIUM (2)
  → Feedback ≥ 0
- LOW (1)
  → Not required

3. Derived output

After applying rules, the system must produce:

a. **Critical dimensions**
   → must meet ≥ +1
b. **Minimum threshold dimensions**
   → must meet ≥ 0
c. **Non-critical dimensions**
   → supporting only

Enforcement rules:

Rule 1; must pass enforcement before decision.

If a **critical dimension= 0**

→ **Decision: Insufficient Evidence (No Judgement)**
→ **Action: Reassess Required**

Constraint:

Cannot assign:

- **Proceed**
- **Proceed with Conditions**
- **Do Not Proceed**

Rationale:

A critical requirement has not been evaluated. Any decision would be invalid.

Rule 2; no compensation.

Strength in **non-critical dimensions** cannot offset failure in **critical dimensions**.

Application:

- A candidate cannot “make up for” weakness in a required capability
- Prevents over-indexing on isolated strengths

Rule 3; environment specific failure.

Failure in a **critical dimension** = **environmental misfit**, not just weakness.

Application:

- Critical dimensions are defined by environment (via 2k mapping) **excluding Execution**, which is globally enforced
- Failure indicates inability to operate effectively in that context

Rule 4; dynamic thresholding.

Same candidate + different environment = different outcome

Application:

- Thresholds are environment-dependent (via 2h / 2j)
- Decision outcomes must reflect role-specific requirements

Rule 5; critical risk enforcement

Critical dimensions with **negative signals** must directly influence decision stance.

Application:

If a **critical dimension = -1 (with Medium+ confidence)**:

→ Decision must only be:

- **Proceed with Conditions** OR
- **Do Not Proceed**.

If a **critical dimension = -2 (with Medium+ confidence)**:

→ **Do Not Proceed**
→ Immediate decision (no override).

Execution is excluded from environment-based prioritisation and is always treated as a global constraint.

Mapping engine example:

Input (org):

- Structure = No clear process, build as you go
- Pace = Fast, but manageable
- Stakeholder Complexity = Multiple, often conflicting
- Autonomy = Act independently, adjust as needed
- Feedback = Situational

3 options that actually drive hiring decisions in this role:

- Decision-making under ambiguity
- Stakeholder navigation under pressure
- Information gathering & clarification

Maps to:

- Structure: LOW (3)
- Pace: FAST (3)
- Stakeholder Complexity: HIGH (3)
- Autonomy: HIGH (3)
- Feedback: SITUATIONAL (2)

Step 1: Apply mapping rules

- Autonomy (3)
    - → Ownership ≥ +1 (Critical)
    - → Ambiguity ≥ +1 (Critical)
- Pace (3)
    - → Execution ≥ +1 (Critical)
    - → Execution -1 = Do Not Proceed
- Stakeholder (3)
    - → Conflict ≥ +1 (Critical)
- Structure (3)
    - → Ambiguity ≥ +1 (reinforces critical)
- Feedback (2)
    - → Feedback ≥ 0 (Minimum)

Step 2: Output classification

Critical Dimensions:

- Execution (regardless, due to global constraint)
- Ownership
- Ambiguity
- Conflict & Challenge

Minimum Dimensions:

- Feedback & Adaptation

Non-critical Dimensions:

- None in this case

Step 3: Add priority overlay:

Selected:

- Decision-making under ambiguity → Ambiguity
- Stakeholder navigation → Conflict
- Info gathering → Ambiguity

Priority dimensions:

- Ambiguity (rank 1)
- Conflict (rank 2)

Final system state:

- **Critical = 4 dimensions**
- **Minimum = 1**
- **Priority = Ambiguity (primary), Conflict (secondary)**

2j Dimensions Classification Framework

Objective:

Convert environment-driven thresholds (2i) into a clear decision structure by classifying behavioural dimensions into:

- Critical (must-pass)
- Minimum (baseline requirement)
- Non-critical (supporting only)

This classification defines **what determines decision validity vs what influences decision direction**.

1. Critical dimensions

Definition:

Dimensions that are **required for successful operation in the defined environment**.

How they are determined:

A dimension is classified as **Critical** if:

- Mapping engine (2i) sets threshold at **≥ +1**; OR
- Environment defines it as **non-negotiable for success** (e.g. fast pace → Execution)

Rules:

- Must achieve **≥ +1**
- If = 0 → **Insufficient Evidence (No Judgement)**
- If ≤ -1 (with Medium+ confidence) → **Cannot result in Proceed**

Purpose:

Defines decision validity. If a critical dimension is not met, the system cannot confidently assert capability.

2. Minimum dimensions

Definition:

Dimensions required to meet **baseline operating standards**, but not decisive for success.

How they are determined:

A dimension is classified as **Minimum** if:

- Mapping engine (2i) sets threshold at **≥ 0**; AND

- It is not elevated to Critical

Rules:

- Must achieve **≥ 0**
- If ≤ -1 (Medium+ confidence):
    - Contributes to **Risk Stack**
    - Influences decision stance (Proceed → Proceed with Conditions / Do Not Proceed)

Purpose:

Prevents **weak or unstable behavioural patterns** from passing unnoticed.

3. Non-critical dimensions

Definition:

Dimensions that are **not required for success in the defined environment**, but provide additional signals.

Rules:

- No minimum threshold
- Do not block decisions
- Can:
    - Increase confidence
    - Contribute to risk stack (if negative and confident)

Purpose:

Adds **depth, not constraint**.

4. System interaction

Hierarchy of influence:

1. **Critical → determines if decision is valid**
2. **Minimum → determines if decision is safe**
3. **Priority (next section) → determines decision direction**

5. Enforcement rules

- No compensation:
    - Strength in non-critical or minimum dimensions cannot offset failure in critical dimensions
- Classification is **environment-driven only**
    - Not interviewer judgement
    - Not candidate-dependent
- Same candidate + different environment:

    - → Different classification
    - → Different outcome

Summary

1. Critical determines **can they do it**
2. Minimum determines **are they safe**
3. Priority determines **do we back them**

2k. Behavioural Priority Weighting

Objective:

Capture role-specific behavioural priorities and apply them to **borderline decisioning** without altering core pass/fail rules.

Core principle:

Critical dimensions determine validity. Priority dimensions determine directional bias once validity is established.

Input:

Organisation selects and ranks key behaviours (top 2–3):

Examples:

- [**x**] Decision-making under ambiguity
- [**x**] Prioritisation under constraint
- [ ] Stakeholder navigation under pressure
- [**x**] Execution consistency under load
- [ ] Information gathering & clarification

Mapping:

Each selected behaviour must map to an existing dimension:

Behaviour	**Dimension**
Decision-making under ambiguity	Ambiguity
Prioritisation under constraint	Execution
Stakeholder navigation under pressure	Conflict and Challenge
Execution consistency under load	Execution
Information gathering & clarification	Ambiguity

Output:

Produces:

- **Priority Dimensions (set, max 3)**
- Derived from selected behaviours via fixed mapping

Example:

Selected behaviours:
- Execution consistency under load
- Information gathering & clarification

Mapped Priority Dimensions:
- Execution
- Ambiguity

Refer to 2g for **how Priority Dimensions are applied** in practice.

Application rules:

1. No override rule

Priority weighting must not:

- override critical dimension requirements (2k)
- allow compensation for weak or failed dimensions
- change minimum thresholds

2. Borderline influence rule

Applied only when decision = **Proceed with Conditions boundary**

Case A - Positive Lean

If:
- Priority dimension ≥ +1
- AND all critical thresholds met

→ **Proceed with Conditions** may resolve to **Proceed**

Case B - Negative Lean

If:
- Priority dimension = 0
- OR weak confidence

→ **Proceed with Conditions** remains **Proceed with Conditions.**

If:
- Priority dimension ≤ -1

→ **Proceed with Conditions** resolves to **Do Not Proceed**

3. Signal strength requirement

Priority influence only applies when:

- Evidence meets 2d thresholds
- Confidence ≥ Medium

4. Tie-breaking rule

If multiple candidates OR borderline decision:

- Compare highest-ranked priority dimension first
- If equal → move to next
- If still equal → no priority influence applied

(No ordinal ranking output; internal only)

Constraints:

- Maximum 3 priority behaviours
- Must map cleanly to dimensions
- Cannot introduce new scoring logic

System position:

- 2h → defines environment
- 2i → defines requirements
- 2k → adjusts sensitivity
- 2g → produces final decision

Failure conditions:

If priority behaviours:

- Do not map cleanly; or
- Lack supporting evidence

→ **no weighting applied**.

Scoring:

Decision-making under ambiguity:

+2

- Acts decisively without full information
- Frames problems clearly despite uncertainty
- Updates direction quickly when new inputs emerge

+1

- Moves forward with partial clarity but seeks reassurance
- Some hesitation, but progress maintained

-1

- Slows down when information is incomplete
- Over-relies on additional input before acting

-2

- Delays decisions waiting for clarity
- Avoids committing under uncertainty
- Becomes reactive when conditions shift

Prioritisation Under Constraint:

+2

- Re-orders work effectively under competing demands
- Explicitly drops lower-value tasks
- Focuses effort on highest impact outcomes

+1

- Adjusts priorities but lacks full clarity

-1

- Struggles to prioritise when multiple demands exist
- Keeps too many things active

-2

- Cannot make trade-offs
- Continues original plan despite conflicting priorities
- Spreads effort thin with no clear focus

Stakeholder navigation under pressure:

+2

- Engages pushback directly without losing position
- Adapts communication to stakeholder level (user → exec)
- Balances relationship and outcome effectively

+1

- Manages stakeholders adequately but concedes ground too early

-1

- Avoids tension or difficult conversations
- Struggles with senior or influential stakeholders

-2

- Concedes under pressure without justification
- Prioritises agreement over outcome
- Escalates unnecessarily instead of engaging directly

Execution consistency under load:

+2
- Maintains output quality under pressure
- Delivers consistently despite competing demands
- Keeps momentum without sacrificing structure

+1
- Output mostly consistent, minor drop under pressure

-1
- Output becomes inconsistent under load
- Misses details or follow-through

-2
- Breaks down under pressure
- Loses track of commitments
- Work becomes reactive or incomplete

Information gathering and clarification

+2
- Identifies critical gaps quickly
- Asks targeted, high-value questions
- Clarifies before acting without slowing momentum

+1
- Gathers sufficient info but not always efficiently

-1
- Misses important context
- Asks broad or unfocused questions

-2
- Proceeds with flawed or missing understanding
- Fails to identify key risks or dependencies
- Relies on assumptions without validation

2I. Decision Framing Layer

Objective:

Determine how decisions will be presented and with what critical information.

Output structure:

Every decision must output:

- Decision [stance]
- Why [strengths]
    - Top validated behavioural signals
- Risks
    - Dimension + confidence tagged
- Environment fit
    - Thrive / Operate / Struggle
- Trade-off
    - “Proceeding due to X despite Y”
- Conditions
    - Mandatory if “Proceed with Conditions”

Conditions layer:

If → **Proceed with Conditions** the system **must** output **specific mitigations tied to risks**.

Example:

- Risk: Execution (-1, Medium)
- → Condition: structured delivery checkpoints

Override capture:

If a user overrides they must select **what drove the override**?

Options:

- Communication style
- Energy / presence
- Experience intuition
- Cultural alignment
- Other (forced input/text)

Outcome:

Stored as **unstructured signal → converted into trackable data.**

Decision: Proceed with conditions

Why:
- Strong collaboration and communication signals
- Comfortable and capable in dealing with ambiguity

Risks:
- Execution ownership (Med confidence)
- Stakeholder management depth (Low confidence)

Environment fit:
- Likely to struggle in high-autonomy, fast-paced settings

Trade-off being made:
- Hiring for collaboration strength despite execution risk

Suggested mitigation:

- Tight performance checkpoints (0–60 days)
- Clear ownership boundaries

3. Decision Routines

Notes:

1. Probe logic and signal extraction must conform to LLM prompt structure defined in 4g
2. Interview include coverage similar to:
    a. Constraint
        i. Time
        ii. Resources
        iii. Conflicting goals
    b. 2. Irreversibility
        i. “You can’t do both”
        ii. “You’ve already committed to X”
    c. 3. Follow-up pressure
        i. “Why that over the other?”
        ii. “What are you dropping?”
        iii. “What happens if that fails?”

–

4. Decision Architecture

Notes:

–

4f. Archetype Layer: Contextual Intelligence & Pattern Learning

Role Archetype Model:

Defines: How roles/environments are structured and classified

Includes:

- Company stage (early / growth / scaled)
- Business velocity (stable / high-growth / hypergrowth)
- Operating model (structured vs fluid)
- Risk tolerance
- Decision-making style (centralised vs distributed)
- Functional context (e.g. QA in fintech doesn't equal QA in gov)

Candidate Archetype Model:

Defines: How candidates are grouped based on behavioural evidence patterns

Built from:

- 5D behavioural scoring outputs
- Signal combinations
- Contextual performance indicators

Examples (internally, firstly):

- "High ambiguity / low structure executor"
- "Structured environment optimiser"
- "High ownership / low challenge risk profile"

NOT personality typing. It's **observed operating patterns.

Archetype Mapping Layer:

Defines: How candidate archetypes are matched to role archetypes.

Includes:

- Success pattern detection - which behavioural combinations succeed in early-stage fintech QA vs scaled SaaS QA
- Failure pattern detection - which profiles fail despite looking strong on paper
- Behaviour-environment compatibility scoring - which signals are predictive vs noise

Learning Loop and Evolution:

Defines: How archetypes improve over time via system data

Includes:

- Pattern reinforcement
- Drift detection (tie to 5d)
- Reclassification logic

Monestisation oppy:

- Predictive hiring benchmarks
- Org diagnostics (“you say you want X, but your environment selects for Y”)
- Talent supply mapping by archetype
- Premium matching based on **probability of success, not similarity**
- Profiling datasets for external projects

–

4g. LLM Execution Layer: Prompting & Scoring Enforcement

Prompt Architecture: (hard constraints)

System Prompt - non negotiable rules. This is the “constitution.” It never changes per candidate.

“

You are an evaluation engine within the Talenti Decision System.

You must:
- Evaluate ONLY based on observed evidence
- Ignore tone, confidence, and articulation quality
- Reject any claim not supported by action + decision + context
- Score strictly according to defined behavioural criteria
- Default to 0 if evidence is insufficient

You must NOT:
- Infer intent without evidence
- Fill gaps with assumptions
- Reward hypothetical or theoretical answers
- Override scoring rules

If evidence is incomplete:
- Mark as insufficient
- Do not upgrade score

All outputs must follow the defined JSON schema exactly.

“

Dimension Scoring Prompt:

Each dimension gets its own evaluation block.

Dimension Evaluation Template: (example: Ambiguity)

“

Evaluate the candidate on the dimension: Ambiguity

Definition:
Ability to make effective decisions under incomplete information, unclear goals, and uncertain conditions.

You must assess:

1. Did the candidate act without full clarity?
2. Did the candidate make decisions with incomplete data?
3. Did the candidate demonstrate reasoning and trade-offs?
4. Did the candidate adapt when new information emerged?

Evidence Requirements:
- Context must include ambiguity (unclear goal, missing info, or uncertainty)
- Must include a clear decision made by the candidate
- Must include reasoning behind that decision

Scoring Rules:
+2 → Multiple examples, clear trade-offs, adaptation shown
+1 → One solid example with decision and reasoning
0 → No real decision under ambiguity
-1 → Avoided decision or required excessive direction
-2 → Failed under ambiguity or acted recklessly without reasoning

Output:
- Score: [-2, -1, 0, +1, +2]
- Evidence summary (max 3 bullet points)
- Confidence: [High, Medium, Low]
- Evidence gaps (if any)

“

Replicated across the 5 dimensions - no variation, no interpretation drift.

Scoring Schema:

Structured output schema (JSON only)

No free text allowed outside of this.

“

```
{
  "candidate_id": "string",
  "dimension_scores": {
    "ownership": {
      "score": 1,
      "confidence": "medium",
      "evidence": ["..."],
      "gaps": ["..."]
    },
    "execution": {
```

```
      "score": 0,
      "confidence": "low",
      "evidence": [],
      "gaps": ["no measurable outcome"]
    },
    "ambiguity": {
      "score": 1,
      "confidence": "high",
      "evidence": ["..."],
      "gaps": []
    },
    "conflict_challenge": {
      "score": -1,
      "confidence": "medium",
      "evidence": ["avoided disagreement with manager"],
      "gaps": []
    },
    "feedback_adaptation": {
      "score": 0,
      "confidence": "low",
      "evidence": [],
      "gaps": ["no behavioural change observed"]
    }
  },
  "decision": {
    "outcome": "Do Not Proceed",
    "reason": "Conflict & Challenge ≤ -1 (kill condition triggered)",
    "confidence": "high"
  },
  "integrity": {
    "status": "clean",
    "notes": []
  }
}
```


“

Decision Logic Enforcement:

Decision engine rules - must be coded, not interpreted.

LLM does not decide freely, it applies **rules** like:

“

IF any dimension ≤ -1 AND dimension is critical
→ **Do Not Proceed**

IF fewer than 3 dimensions ≥ +1
→ **Proceed with Conditions** or **Do Not Proceed**

IF no +2 present
→ Cannot Proceed

IF conflicting signals exist
→ Default Hold unless severity rule applies

“

These rules must be **explicitly injected into the prompt OR handled post-LLM deterministically** (preferrably both).

Anti Gaming Enforcement:

Signal validation prompts - added with a second pass.

“

Check for:
- Rehearsed answers
- Generic frameworks
- Lack of specificity
- Hypothetical disguised as real

If detected:
→ Invalidate signal
→ Reduce confidence
→ Flag integrity

“

**Don’t do this

- Let LLM “summarise and score” in one go → too loose
- Allowing free-text reasoning → introduces drift
- Not enforcing schema → breaks consistency
- Mixing dimensions in one prompt → muddies scoring

5. Decision Governance

Notes:

–

6. Cross-Cutting Controls: Signal Integrity

Notes:

–

6a.

6g. Skills Signal Validation Rules

- Reject theoretical answers
- Reject templated outputs
- Reject constraint avoidance
- Reject partial task completion

–

7. Decision Contract

Notes:

1. All outputs must be generated via structured LLM schema (4g), no free-text decisions permitted
2. Core output pattern: (behavioural summary layer)
    a. Under high ambiguity + high stakeholder pressure:
        i. Maintains decision clarity
        ii. Re-prioritises effectively when disrupted
        iii. Engages resistance directly without escalation
        iv. Sustains forward momentum under constraint
3. Primary output for decision card:
    a. Decision
    b. Why
    c. Environment Fit
    d. Key Risk
    e. Trade-off
    f. Mitigation

–

7g. Candidate Feedback Layer

Principle:

Candidate feedback must be **derived from the same system outputs**, not reinterpreted.

2 tiers:

1. Standard Feedback (default, all candidates)

Objective:

Provide useful but controlled insight without:

- Over-signalling fit
- Creating legal risk
- Undermining hiring decisions

Content structure:

- Behavioural snapshot (light)
    - 1–2 strengths (non-ranked, non-scored language)
    - 1–2 development areas
- General guidance
    - High-level behavioural advice (not tied to the specific org decision)
- **No environment mapping**
- **No explicit scores**
- **No decision justification**

Tone:

- Neutral
- Developmental
- Non-committal

2. Enhanced Feedback (paid / premium / subscription orgs)

Objective:

Turn feedback into **candidate value + system differentiatio**n

Content structure:

1. Behavioural Profile
    a. Top strengths (mapped to dimensions)
    b. Key risks (clear, but framed constructively)
2. Environment Fit Mapping

        a. “You are most likely to perform well in environments that are:”
            i. High/Low structure
            ii. High/Low ambiguity
            iii. Etc.
        b. AND
            i. Concrete examples: i.e.,
                1. Early-stage / scale-up / enterprise
                2. Fast-moving vs stable orgs
3. Behavioural patterns
    a. Observed tendencies:
        i. e.g. “Strong executor in structured contexts, less effective under ambiguity”
4. Development Pathways
    a. Specific improvement areas tied to:
        i. Behaviours, not personality

What we **will not** do/uncover:

- Final hiring decision reasoning in full
- Org-specific requirements (“they wanted X, you lacked Y”)
- Scores (-2 to +2)
- Confidence levels

System integration

Inputs:

- Dimension scores
- Confidence levels
- Risk stack
- Environment profile
- Mapping outputs (critical/minimum/priority)

Transformation rules

Standard feedback:

Strip:

- scores
- confidence
- environment specificity

Retain:

- generalised behavioural signals

Enhanced feedback:

Translate:

- dimensions → behavioural traits
- environment → “fit zones”

Apply:

- abstraction layer (no raw system language)

Constraints:

- Must not contradict final decision
- Must not imply hiring outcome
- Must not expose internal thresholds